Teaching Bayes: The Essential Parts

Rebecca C. Steorts
ISBA World Meeting 2016

Lecture 2: Intro to Gibbs Sampling

Intro to Markov chain Monte Carlo (MCMC)

Goal: sample from $f(x)$, or approximate $E_f[h(X)]$.

Function $f(x)$ is very complicated and hard to sample from.

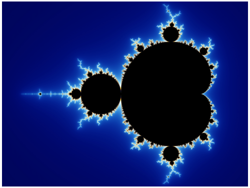

How to deal with this?

1. What's a simple way?
2. What are two other ways?
3. What happens in high dimensions?

High dimensional spaces

- In low dimensions, importance and rejection sampling work pretty well.
- But in high dimensions, a proposal $g(x)$ that worked in 2-D, often doesn't mean that it will work in any dimension.
- Why? It's hard to capture high dimensional spaces!

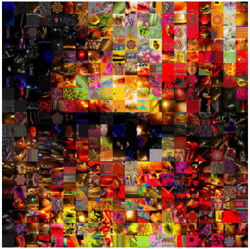

Figure 1: A high dimensional space (many images).

We turn to Markov chain Monte Carlo (MCMC).

Intution

Imagine that we have a complicated function f below and it's high probability regions are represented in green.

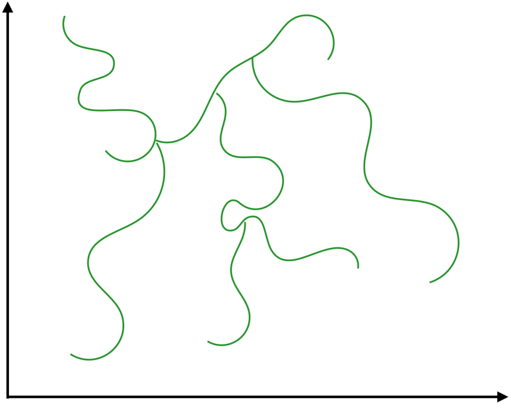

Figure 2: Example of a Markov chain

Intution

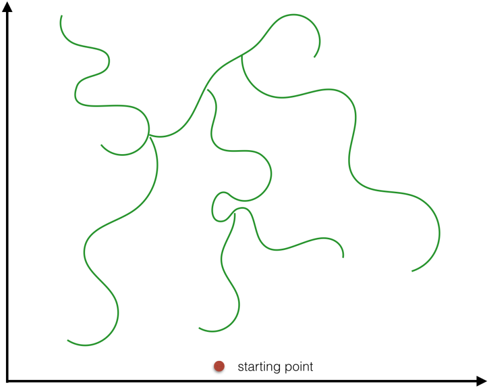

Figure 3: Example of a Markov chain and red starting point

Intution

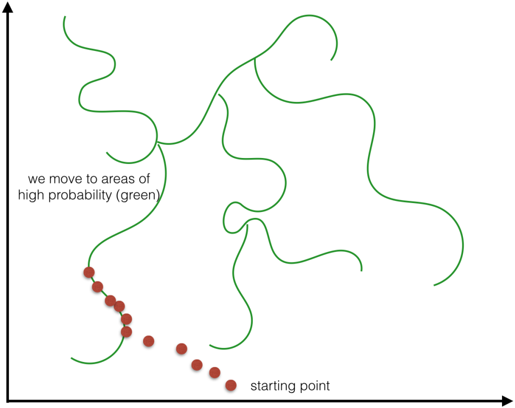

Figure 4: Example of a Markov chain and moving from the starting point to a high probability region.

What is Markov Chain Monte Carlo

- Markov Chain – where we go next only depends on our last state (the Markov property).
- Monte Carlo – just simulating data.

Why MCMC?

(a) the region of high probability tends to be “connected”

  - That is, we can get from one point to another without going through a low-probability region, and

(b) we tend to be interested in the expectations of functions that are relatively smooth and have lots of “symmetries”

  - That is, one only needs to evaluate them at a small number of representative points in order to get the general picture.

Advantages/Disadvantages of MCMC:

Advantages:

- applicable even when we can't directly draw samples
- works for complicated distributions in high-dimensional spaces, even when we don't know where the regions of high probability are
- relatively easy to implement
- fairly reliable

Disadvantages:

- slower than simple Monte Carlo or importance sampling (i.e., requires more samples for the same level of accuracy)
- can be very difficult to assess accuracy and evaluate convergence, even empirically

Two-stage Gibbs sampler

- Suppose $p(x, y)$ is a p.d.f. or p.m.f. that is difficult to sample from directly.
- Suppose, though, that we *can* easily sample from the conditional distributions $p(x|y)$ and $p(y|x)$.
- The Gibbs sampler proceeds as follows:
    1. set x and y to some initial starting values
    2. then sample $x|y$, then sample $y|x$,
       then $x|y$, and so on.

Two-stage Gibbs sampler

0. Set (x_0, y_0) to some starting value.
1. Sample $x_1 \sim p(x|y_0)$, that is, from the conditional distribution $X \mid Y = y_0$.

   Current state: (x_1, y_1)

   Sample $y_1 \sim p(y|x_1)$, that is, from the conditional distribution $Y \mid X = x_1$.

   Current state: (x_1, y_1)
2. Sample $x_2 \sim p(x|y_1)$, that is, from the conditional distribution $X \mid Y = y_1$.

   Current state: (x_2, y_1)

   Sample $y_2 \sim p(y|x_2)$, that is, from the conditional distribution $Y \mid X = x_2$.

   Current state: (x_2, y_2)

   $\vdots$

Repeat iterations 1 and 2, M times.

This procedure defines a sequence of pairs of random variables

$$(X_0, Y_0), (X_1, Y_1), (X_2, Y_2), (X_3, Y_3), \ldots$$

Markov chain and dependence

$$(X_0, Y_0), (X_1, Y_1), (X_2, Y_2), (X_3, Y_3), \ldots$$

satisfies the property of being a Markov chain.

The conditional distribution of (X_i, Y_i) given all of the previous pairs depends only on (X_{i-1}, Y_{i-1})

$(X_0, Y_0), (X_1, Y_1), (X_2, Y_2), (X_3, Y_3), \ldots$ are not iid samples (Think about why).

Ideal Properties of MCMC

- (x_0, y_0) chosen to be in a region of high probability under $p(x, y)$, but often this is not so easy.
- We run the chain for M iterations and discard the first B samples $(X_1, Y_1), \ldots, (X_B, Y_B)$. This is called *burn-in*.
- Typically: if you run the chain long enough, the choice of B doesn't matter.
- Roughly speaking, the performance of an MCMC algorithm—that is, how quickly the sample averages $\frac{1}{N}\sum_{i=1}^{N} h(X_i, Y_i)$ converge—is referred to as the *mixing rate*.
- An algorithm with good performance is said to "have good mixing", or "mix well".

Toy Example

Suppose we want to sample from the bivariate distribution:

$$p(x, y) \propto e^{-xy} \mathbb{1}(x, y \in (0, c))$$

where $c > 0$, and $(0, c)$ denotes the (open) interval between 0 and c. (This example is due to Casella & George, 1992.)

Toy Example

- The Gibbs sampling approach is to alternately sample from $p(x|y)$ and $p(y|x)$.
- Note $p(x, y)$ is symmetric with respect to x and y.
- Hence, only need to derive one of these and then we can get the other one by just swapping x and y.
- Let's look at $p(x|y)$.

Toy Example

$$p(x, y) \propto e^{-xy} \mathbb{1}(x, y \in (0, c))$$

$$p(x|y) \underset{x}{\propto} p(x, y) \underset{x}{\propto} e^{-xy} \mathbb{1}(0 < x < c) \underset{x}{\propto} \mathrm{Exp}(x|y) \mathbb{1}(x < c).$$[1]

- $p(x|y)$ is a *truncated* version of the $\mathrm{Exp}(y)$ distribution
- It is the same as taking $X \sim \mathrm{Exp}(y)$ and conditioning on it being less than c, i.e., $X \mid X < c$.
- Let's refer to this as the $\mathrm{TExp}(y, (0, c))$ distribution.

[1] Under $\propto$, we write the random variable (x) for clarity.

Toy Example

An easy way to generate a sample from $Z \sim \mathrm{TExp}(\theta, (0, c))$, is:

1. Sample $U \sim \mathrm{Uniform}(0, F(c|\theta))$ where

$$F(x|\theta) = 1 - e^{-\theta x}$$

   is the $\mathrm{Exp}(\theta)$ c.d.f.

2. Set $Z = F^{-1}(U|\theta)$ where

$$F^{-1}(u|\theta) = -(1/\theta)\log(1-u)$$

   is the inverse c.d.f. for $u \in (0, 1)$.

Verify the last step on your own.

Let's apply Gibbs sampling, denoting $S = (0, c)$.

0. Initialize $x_0, y_0 \in S$.
1. Sample $x_1 \sim \mathrm{TExp}(y_0, S)$, then sample $y_1 \sim \mathrm{TExp}(x_1, S)$.
2. Sample $x_2 \sim \mathrm{TExp}(y_1, S)$, then sample $y_2 \sim \mathrm{TExp}(x_2, S)$.

   $\vdots$

N. Sample $x_N \sim \mathrm{TExp}(y_{N-1}, S)$, sample $y_N \sim \mathrm{TExp}(x_N, S)$.

Figure 5 demonstrates the algorithm, with $c = 2$ and initial point $(x_0, y_0) = (1, 1)$.

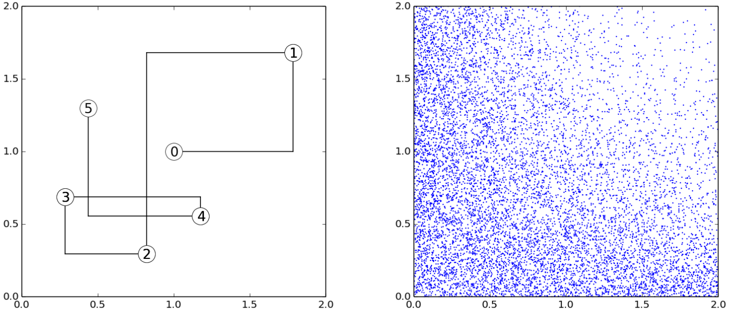

Figure 5: (Left) Schematic representation of the first 5 Gibbs sampling iterations/sweeps/scans. (Right) Scatterplot of samples from 10^4 Gibbs sampling iterations.

Pareto example

Distributions of sizes and frequencies often tend to follow a “power law” distribution.

- wealth of individuals
- size of oil reserves
- size of cities
- word frequency
- returns on stocks

Power Law Distribution

The Pareto distribution with shape $\alpha > 0$ and scale $c > 0$ has p.d.f.

$$\mathrm{Pareto}(x|\alpha, c) = \frac{\alpha c^{\alpha}}{x^{\alpha+1}} \mathbb{1}(x > c) \propto \frac{1}{x^{\alpha+1}} \mathbb{1}(x > c).$$

This is referred to as a power law distribution, because the p.d.f. is proportional to x raised to a power. Notice that c is a lower bound on the observed values. In this example, we'll see how Gibbs sampling can be used to perform inference for α and c.

Rank	City	Population
1	Charlotte	731424
2	Raleigh	403892
3	Greensboro	269666
4	Durham	228330
5	Winston-Salem	229618
6	Fayetteville	200564
7	Cary	135234
8	Wilmington	106476
9	High Point	104371
10	Greenville	84554
11	Asheville	85712
12	Concord	79066
⋮	⋮	⋮
44	Havelock	20735
45	Carrboro	19582
46	Shelby	20323
47	Clemmons	18627
48	Lexington	18931
49	Elizabeth City	18683
50	Boone	17122

Parameter Interpretations

- α tells us the scaling relationship between the size of cities and their probability of occurring.
  - Let $\alpha = 1$.
  - Density looks like $1/x^{\alpha+1} = 1/x^2$.
  - Cities with 10,000–20,000 inhabitants occur roughly $10^{\alpha+1} = 100$ times as frequently as cities with 100,000–110,000 inhabitants.
- c represents the cutoff point—any cities smaller than this were not included in the dataset.

To keep things as simple as possible, let's use an (improper) default prior:

$$p(\alpha, c) \propto \mathbb{1}(\alpha, c > 0).$$

Recall from Module 4:

- An *improper/default prior* is a nonnegative function of the parameters which integrates to infinity.
- Often (but not always!) the resulting "posterior" will be proper.
- It is important that the "posterior" be proper, since otherwise the whole Bayesian framework breaks down.

Recall

$$p(x|\alpha, c) = \frac{\alpha c^{\alpha}}{x^{\alpha+1}} \mathbb{1}(x > c) \tag{1}$$

$$\mathbb{1}(\alpha, c > 0) \tag{2}$$

Let's derive the posterior:

$$\begin{aligned}
p(\alpha, c|x_{1:n}) &\underset{\alpha,c}{\overset{\text{def}}{\propto}} p(x_{1:n}|\alpha, c)p(\alpha, c) \\
&\underset{\alpha,c}{\propto} \mathbb{1}(\alpha, c > 0) \prod_{i=1}^{n} \frac{\alpha c^{\alpha}}{x_i^{\alpha+1}} \mathbb{1}(x_i > c) \\
&= \frac{\alpha^n c^{n\alpha}}{(\prod x_i)^{\alpha+1}} \mathbb{1}(c < x_*) \mathbb{1}(\alpha, c > 0) \qquad (3)
\end{aligned}$$

where $x_* = \min\{x_1, \ldots, x_n\}$.

As a joint distribution on (α, c),

- this does not seem to have a recognizable form,
- and it is not clear how we might sample from it directly.

Let's try Gibbs sampling!

To use Gibbs, we need to be able to sample $\alpha|c, x_{1:n}$ and $c|\alpha, x_{1:n}$.

By Equation 3, we find that

$$\begin{aligned} p(\alpha|c, x_{1:n}) \underset{\alpha}{\propto} p(\alpha, c|x_{1:n}) \underset{\alpha}{\propto} \frac{\alpha^n c^{n\alpha}}{(\prod x_i)^\alpha} \mathbb{1}(\alpha > 0) \\ = \alpha^n \exp\big(-\alpha(\textstyle\sum \log x_i - n \log c)\big) \mathbb{1}(\alpha > 0) \\ \underset{\alpha}{\propto} \mathrm{Gamma}\big(\alpha \,\big|\, n+1,\, \textstyle\sum \log x_i - n \log c\big), \end{aligned}$$

and

$$p(c|\alpha, x_{1:n}) \underset{c}{\propto} p(\alpha, c|x_{1:n}) \underset{c}{\propto} c^{n\alpha} \mathbb{1}(0 < c < x_*),$$

which we will define to be Mono(α, x_*)

Defining the Mono distribution

For $a > 0$ and $b > 0$, define the distribution $\text{Mono}(a, b)$ (for monomial) with p.d.f.

$$\text{Mono}(x|a, b) \propto x^{a-1}\mathbb{1}(0 < x < b).$$

Since $\int_0^b x^{a-1}dx = b^a/a$, we have

$$\text{Mono}(x|a, b) = \frac{a}{b^a}x^{a-1}\mathbb{1}(0 < x < b),$$

and for $0 < x < b$, the c.d.f. is

$$F(x|a, b) = \int_0^x \text{Mono}(y|a, b)dy = \frac{a}{b^a}\frac{x^a}{a} = \frac{x^a}{b^a}.$$

To use the inverse c.d.f. technique, we solve for the inverse of F on $0 < x < b$: Let $u = \frac{x^a}{b^a}$ and solve for x.

$$u = \frac{x^a}{b^a} \tag{4}$$

$$b^a u = x^a \tag{5}$$

$$b u^{1/a} = x \tag{6}$$

Can sample from $\mathrm{Mono}(a, b)$ by drawing $U \sim \mathrm{Uniform}(0, 1)$ and setting $X = bU^{1/a}$.[2]

[2] It turns out that this is an inverse of the Pareto distribution, in the sense that if $X \sim \mathrm{Pareto}(\alpha, c)$ then $1/X \sim \mathrm{Mono}(\alpha, 1/c)$.

So, in order to use the Gibbs sampling algorithm to sample from the posterior $p(\alpha, c|x_{1:n})$, we initialize α and c, and then alternately update them by sampling:

$$\begin{aligned}
\alpha|c, x_{1:n} &\sim \text{Gamma}\left(n+1,\, \textstyle\sum \log x_i - n \log c\right) \\
c|\alpha, x_{1:n} &\sim \text{Mono}(n\alpha + 1,\, x_*).
\end{aligned}$$

Ways of visualizing results

Traceplots. A traceplot simply shows the sequence of samples, for instance $\alpha_1, \ldots, \alpha_N$, or $c_1, \ldots, c_N$. Traceplots are a simple but very useful way to visualize how the sampler is behaving.

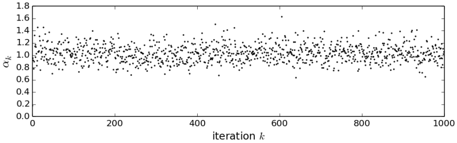

Figure 6: Traceplot of α

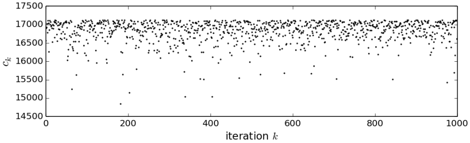

Figure 7: Traceplot of c.

Estimated density. We are primarily interested in the posterior on α, since it tells us the scaling relationship between the size of cities and their probability of occurring.

By making a histogram of the samples $\alpha_1, \ldots, \alpha_N$, we can estimate the posterior density $p(\alpha|x_{1:n})$.

The two vertical lines indicate the lower ℓ and upper u boundaries of an (approximate) 90% credible interval $[\ell, u]$—that is, an interval that contains 90% of the posterior probability:

$$\mathbb{P}\big(\boldsymbol{\alpha} \in [\ell, u] \big| x_{1:n}\big) = 0.9.$$

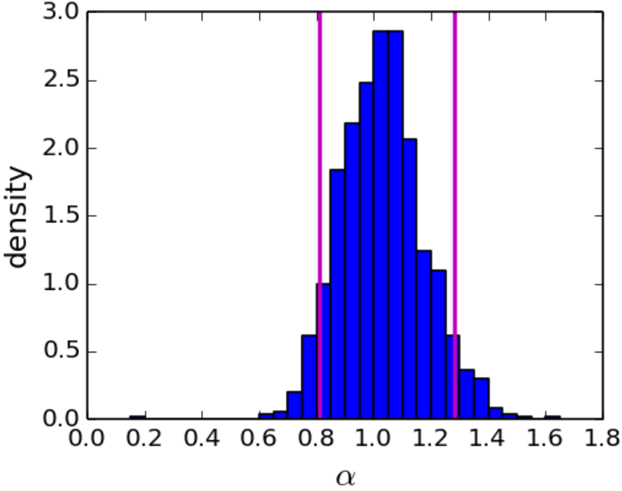

Figure 8: Estimated density of $\alpha|x_{1:n}$ with $\approx$ 90 percent credible intervals.

Running averages. Panel (d) shows the running average $\frac{1}{k}\sum_{i=1}^{k}\alpha_i$ for $k = 1, \ldots, N$.

In addition to traceplots, running averages such as this are a useful heuristic for visually assessing the convergence of the Markov chain.

The running average shown in this example still seems to be meandering about a bit, suggesting that the sampler needs to be run longer (but this would depend on the level of accuracy desired).

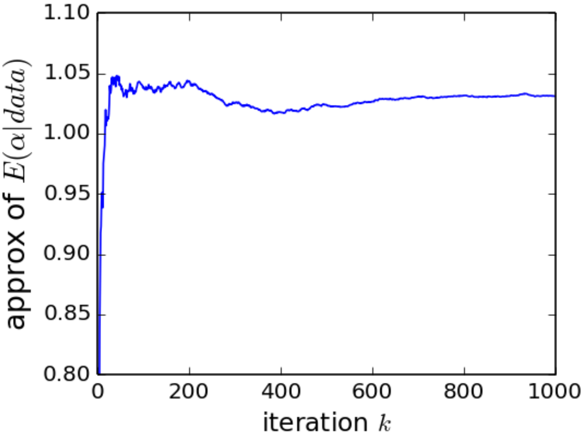

Figure 9: Running average plot

Survival function

A survival function is defined to be

$$S(x) = \mathbb{P}(X > x) = 1 - \mathbb{P}(X \leq x).$$

Power law distributions are often displayed by plotting their survival function $S(x)$, on a log-log plot.

Why? $S(x) = (c/x)^\alpha$ for the $\mathrm{Pareto}(\alpha, c)$ distribution and on a log-log plot this appears as a line with slope $-\alpha$.

The posterior survival function (or more precisely, the posterior predictive survival function), is $S(x|x_{1:n}) = \mathbb{P}(X_{n+1} > x \mid x_{1:n})$.

Figure 10(e) shows an empirical estimate of the survival function (based on the empirical c.d.f., $\hat{F}(x) = \frac{1}{n}\sum_{i=1}^{n} \mathbb{1}(x \geq x_i)$) along with the posterior survival function, approximated by

$$\begin{aligned} S(x|x_{1:n}) &= \mathbb{P}(X_{n+1} > x \mid x_{1:n}) && (7) \\ &= \int \mathbb{P}(X_{n+1} > x \mid \alpha, c) p(\alpha, c|x_{1:n}) d\alpha dc && (8) \\ &\approx \frac{1}{N}\sum_{i=1}^{N} \mathbb{P}(X_{n+1} > x \mid \alpha_i, c_i) = \frac{1}{N}\sum_{i=1}^{N} (c_i/x)^{\alpha_i}. && (9) \end{aligned}$$

This is computed for each x in a grid of values.

[Think about why each line is true on your own].

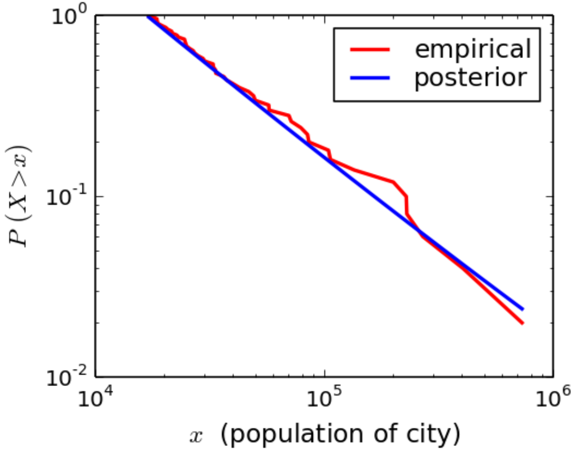

Figure 10: Empirical vs posterior survival function

Questions you should be able to answer!

- When should we use MCMC in a Bayesian setting?
- When would we use an MCMC over Importance sampling and Rejection sampling?
- What is a Gibbs sampler?
- What are simple diagnostics of MCMC?
- Are we guaranteed convergence of the Markov chain emprically?
- What do are diagnostics really tell us?