

# Bayesian Statistics in R - STA 360/601

Abbas Zaidi and Rebecca C. Steorts

## 1 Agenda

In class, you saw the Beta-Bernoulli model. We will now use this to solve a very real problem! Suppose I wish to determine whether the probability that a worker will fake an illness is truly 1%. Your task is to assist me!

1. Simulate some data using the `rbinom` function of size  $n = 100$  and probability equal to 1%. Remember to `set.seed(123)` so that you can replicate your results.
2. Write a function that takes as its inputs that data you simulated (or any data of the same type) and a sequence of  $\theta$  values of length 1000 and produces Likelihood values based on the Binomial Likelihood. Plot your sequence and its corresponding Likelihood function.
3. Write a function that takes as its inputs prior parameters `a` and `b` for the Beta-Bernoulli model and the observed data, and produces the posterior parameters you need for the model. Generate the posterior parameters for a non-informative prior i.e.  $(a,b) = (1,1)$  and for an informative case  $(a,b) = (3,1)$
4. Create two plots, one for the informative and one for the non-informative case to show the posterior distribution and superimpose the prior distributions on each along with the likelihood. What do you see? Remember to turn the y-axis ticks off since superimposing may make the scale non-sense.