Module 8: Introduction to Gibbs Sampling

Rebecca C. Steorts

Agenda

- Gibbs sampling (two-stage sampler)
- Exponential example
- Normal example
- Pareto example
- Diagnostics
- Multi-stage sampler
- Missing Data

Gibbs sampler

- Suppose $p(x, y)$ is a p.d.f. or p.m.f. that is difficult to sample from directly.
- Suppose, though, that we *can* easily sample from the conditional distributions $p(x|y)$ and $p(y|x)$.
- The Gibbs sampler proceeds as follows:
  1. set x and y to some initial starting values
  2. then sample $x|y$, then sample $y|x$, then $x|y$, and so on.

Gibbs sampler

0. Set (x_0, y_0) to some starting value.
1. Sample $x_1 \sim p(x|y_0)$, that is, from the conditional distribution $X \mid Y = y_0$.
   Current state: (x_1, y_0)
   Sample $y_1 \sim p(y|x_1)$, that is, from the conditional distribution $Y \mid X = x_1$.
   Current state: (x_1, y_1)
2. Sample $x_2 \sim p(x|y_1)$, that is, from the conditional distribution $X \mid Y = y_1$.
   Current state: (x_2, y_1)
   Sample $y_2 \sim p(y|x_2)$, that is, from the conditional distribution $Y \mid X = x_2$.
   Current state: (x_2, y_2)
   $\vdots$

Repeat iterations 1 and 2, M times.

Gibbs sampler

This procedure defines a sequence of pairs of random variables

$$(X_0, Y_0), (X_1, Y_1), (X_2, Y_2), (X_3, Y_3), \ldots$$

Markov chain and dependence

$$(X_0, Y_0), (X_1, Y_1), (X_2, Y_2), (X_3, Y_3), \ldots$$

satisfies the property of being a Markov chain.

The conditional distribution of (X_{i+1}, Y_{i+1}) given all of the previous pairs depends only on (X_i, Y_i)

Example: The conditional distribution of (X_5, Y_5) given all of the previous pairs depends only on (X_4, Y_4)

$(X_0, Y_0), (X_1, Y_1), (X_2, Y_2), (X_3, Y_3), \ldots$ are not iid samples (Think about why).

Ideal Properties of MCMC

- (x_0, y_0) chosen to be in a region of high probability under $p(x, y)$, but often this is not so easy.
- We run the chain for M iterations and discard the first B samples $(X_1, Y_1), \ldots, (X_B, Y_B)$. This is called *burn-in*.
- Typically: if you run the chain long enough, the choice of B doesn't matter.
- Roughly speaking, the performance of an MCMC algorithm—that is, how quickly the sample averages $\frac{1}{N}\sum_{i=1}^{N} h(X_i, Y_i)$ converge—is referred to as the *mixing rate*.
- An algorithm with good performance is said to "have good mixing", or "mix well".

Exponential Example

Consider the following Exponential model for observation(s) $x = (x_1, \ldots, x_n)$.[1]:

$$p(x|a, b) = ab \exp(-abx) I(x > 0)$$

and suppose the prior is

$$p(a, b) = \exp(-a - b) I(a, b > 0).$$

You want to sample from the posterior $p(a, b|x)$.

[1] Please note that in the attached data there are 40 observations, which can be found in data-exponential.csv.

Conditional distributions

$$
\begin{aligned}
p(\boldsymbol{x}|a,b) &= \prod_{i=1}^{n} p(x_i|a,b) \\
&= \prod_{i=1}^{n} ab\exp(-abx_i) \\
&= (ab)^n \exp\left(-ab\sum_{i=1}^{n} x_i\right).
\end{aligned}
$$

The function is symmetric for a and b, so we only need to derive $p(a|\boldsymbol{x},b)$.

Conditional distributions

This conditional distribution satisfies

$$\begin{aligned} p(a|\boldsymbol{x}, b) &\propto_a p(a, b, \boldsymbol{x}) \\ &= p(\boldsymbol{x}|a, b)p(a, b) \\ &= \text{fill in full details for lab this week} \end{aligned}$$

Gibbs sampling code

```
knitr::opts_chunk$set(cache=TRUE)
library(MASS)
data <- read.csv("data-exponential.csv", header = FALSE)
```

Gibbs sampling code

```
##########################################
# This function is a Gibbs sampler
#
# Args
#   start.a: initial value for a
#   start.b: initial value for b
#   n.sims: number of iterations to run
#   data:  observed data, should be in a
            # data frame with one column
#
# Returns:
#   A two column matrix with samples
     #   for a in first column and
# samples for b in second column
##########################################
```

Gibbs sampling code

```
sampleGibbs <- function(start.a, start.b, n.sims, data){
  # get sum, which is sufficient statistic
  x <- sum(data)
  # get n
  n <- nrow(data)
  # create empty matrix, allocate memory for efficiency
  res <- matrix(NA, nrow = n.sims, ncol = 2)
  res[1,] <- c(start.a,start.b)
  for (i in 2:n.sims){
    # sample the values
    res[i,1] <- rgamma(1, shape = n+1,
                       rate = res[i-1,2]*x+1)
    res[i,2] <- rgamma(1, shape = n+1,
                       rate = res[i,1]*x+1)
  }
  return(res)
}
```

Gibbs sampler code

```
# run Gibbs sampler
n.sims <- 10000
# return the result (res)
res <- sampleGibbs(.25,.25,n.sims,data)
head(res)
```

```
##           [,1]      [,2]
## [1,] 0.250000 0.2500000
## [2,] 2.035533 0.1963047
## [3,] 1.968843 0.2652573
## [4,] 1.921762 0.2194616
## [5,] 2.527001 0.1816643
## [6,] 2.497594 0.1717998
```

Exponential Example

You will explore this problem more in lab this week and in your homework.

Toy Example

$$p(x, y) \propto e^{-xy} \mathbb{1}(x, y \in (0, c))$$

$$p(x|y) \underset{x}{\propto} p(x, y) \underset{x}{\propto} e^{-xy} \mathbb{1}(0 < x < c) \underset{x}{\propto} \text{Exp}(x|y) \mathbb{1}(x < c).^2$$

- $p(x|y)$ is a *truncated* version of the $\text{Exp}(y)$ distribution
- It is the same as taking $X \sim \text{Exp}(y)$ and conditioning on it being less than c, i.e., $X \mid X < c$.
- Let's refer to this as the $\text{TExp}(y, (0, c))$ distribution.

[2] Under $\propto$, we write the random variable (x) for clarity.

Toy Example

- The Gibbs sampling approach is to alternately sample from $p(x|y)$ and $p(y|x)$.
- Note $p(x, y)$ is symmetric with respect to x and y.
- Hence, only need to derive one of these and then we can get the other one by just swapping x and y.
- Let's look at $p(x|y)$.

Toy Example

An easy way to generate a sample from $Z \sim \text{TExp}(\theta, (0, c))$, is:

1. Sample $U \sim \text{Uniform}(0, F(c|\theta))$ where

$$F(x|\theta) = 1 - e^{-\theta x}$$

   is the $\text{Exp}(\theta)$ c.d.f.
2. Set $Z = F^{-1}(U|\theta)$ where

$$F^{-1}(u|\theta) = -(1/\theta)\log(1-u)$$

   is the inverse c.d.f. for $u \in (0, 1)$.

Hint: To verify the last step: apply the rejection principle (along with the inverse cdf technique). Verify the last step on your own.

Toy example

Let's apply Gibbs sampling, denoting $S = (0, c)$.

0. Initialize $x_0, y_0 \in S$.
1. Sample $x_1 \sim \mathrm{TExp}(y_0, S)$, then sample $y_1 \sim \mathrm{TExp}(x_1, S)$.
2. Sample $x_2 \sim \mathrm{TExp}(y_1, S)$, then sample $y_2 \sim \mathrm{TExp}(x_2, S)$.

   $\vdots$

N. Sample $x_N \sim \mathrm{TExp}(y_{N-1}, S)$, sample $y_N \sim \mathrm{TExp}(x_N, S)$.

Figure 1 demonstrates the algorithm, with $c = 2$ and initial point $(x_0, y_0) = (1, 1)$.

Toy example

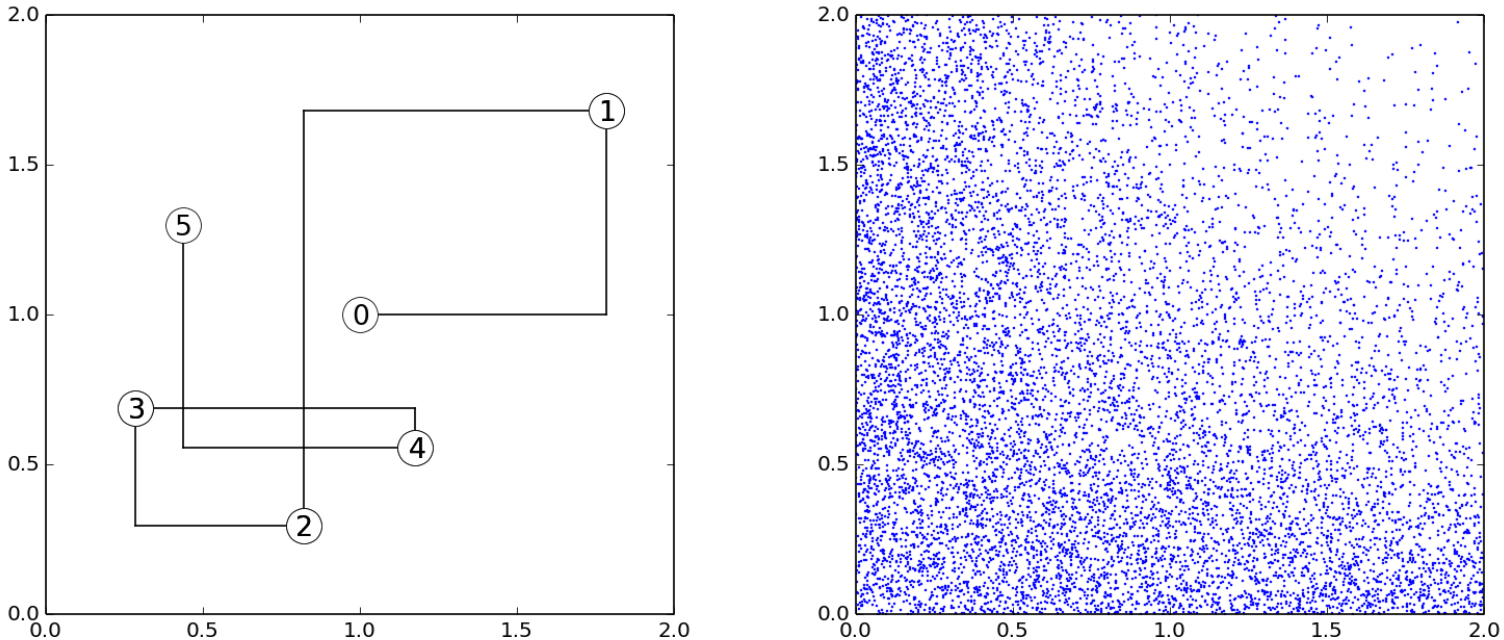

Figure 1: (Left) Schematic representation of the first 5 Gibbs sampling iterations/sweeps/scans. (Right) Scatterplot of samples from 10^4 Gibbs sampling iterations.

Example: Normal with semi-conjugate prior

Consider $X_1, \ldots, X_n | \mu, \lambda \overset{iid}{\sim} \mathcal{N}(\mu, \lambda^{-1})$. Then independently consider

$$\boldsymbol{\mu} \sim \mathcal{N}(\mu_0, \lambda_0^{-1})$$
$$\boldsymbol{\lambda} \sim \text{Gamma}(a, b)$$

This is called a semi-conjugate situation, in the sense that the prior on μ is conjugate for each fixed value of λ, and the prior on λ is conjugate for each fixed value of μ.

For ease of notation, denote the observed data points by $x_{1:n}$.

How does one derive $p(\mu, \lambda \mid x_{1:n})$?

Example

We know that for the Normal–Normal model, we know that for any fixed value of λ,

$$\boldsymbol{\mu}|\lambda, x_{1:n} \sim \mathcal{N}(M_\lambda, L_\lambda^{-1})$$

where

$$L_\lambda = \lambda_0 + n\lambda \text{ and } M_\lambda = \frac{\lambda_0\mu_0 + \lambda\sum_{i=1}^n x_i}{\lambda_0 + n\lambda}.$$

For any fixed value of μ, it is straightforward to derive[3] that

$$\boldsymbol{\lambda}|\mu, x_{1:n} \sim \text{Gamma}(A_\mu, B_\mu) \tag{1}$$

where $A_\mu = a + n/2$ and

$$B_\mu = b + \tfrac{1}{2}\sum(x_i - \mu)^2 = n\hat{\sigma}^2 + n(\bar{x} - \mu)^2$$

where $\hat{\sigma}^2 = \frac{1}{n}\sum(x_i - \bar{x})^2$.

[3] do this on your own

Example

Goal: sample from $p(\mu, \lambda \mid x_{1:n})$

To implement Gibbs sampling in this example, each iteration consists of sampling:

$$\begin{aligned} \boldsymbol{\mu} | \lambda, x_{1:n} &\sim \mathcal{N}(M_\lambda, L_\lambda^{-1}) \\ \boldsymbol{\lambda} | \mu, x_{1:n} &\sim \text{Gamma}(A_\mu, B_\mu). \end{aligned}$$

This will give us samples

$$(\mu_0, \lambda_0), \ldots (\mu_S, \lambda_S)$$

Pareto example

Distributions of sizes and frequencies often tend to follow a "power law" distribution.

- wealth of individuals
- size of oil reserves
- size of cities
- word frequency
- returns on stocks

Power law distribution

The Pareto distribution with shape $\alpha > 0$ and scale $c > 0$ has p.d.f.

$$\text{Pareto}(x|\alpha, c) = \frac{\alpha c^\alpha}{x^{\alpha+1}} \mathbb{1}(x > c) \propto \frac{1}{x^{\alpha+1}} \mathbb{1}(x > c).$$

- This is referred to as a power law distribution, because the p.d.f. is proportional to x raised to a power.
- c is a lower bound on the observed values.
- We will use Gibbs sampling to perform inference for α and c.
- Let X be the population of a city.

Pareto example

Rank	City	Population
1	Charlotte	731424
2	Raleigh	403892
3	Greensboro	269666
4	Durham	228330
5	Winston-Salem	229618
6	Fayetteville	200564
7	Cary	135234
8	Wilmington	106476
9	High Point	104371
10	Greenville	84554
11	Asheville	85712
12	Concord	79066
⋮	⋮	⋮
44	Havelock	20735
45	Carrboro	19582
46	Shelby	20323
47	Clemmons	18627
48	Lexington	18931
49	Elizabeth City	18683

Parameter intepretations

- α tells us the scaling relationship between the size of cities and their probability of occurring.
  - Let $\alpha = 1$.
  - Density looks like $1/x^{\alpha+1} = 1/x^2$.
  - Cities with 10,000–20,000 inhabitants occur roughly $10^{\alpha+1} = 100$ times as frequently as cities with 100,000–110,000 inhabitants (or $10^{\alpha+1}/10 = 10$ times as frequently as cities with 100,000-200,000 inhabitants)
- c represents the cutoff point—any cities smaller than this were not included in the dataset.

Prior selection

For simplicity, let's use an **(improper) default prior**:

$$p(\alpha, c) \propto \mathbb{1}(\alpha, c > 0).$$

Recall:

- An *improper/default prior* is a non-negative function of the parameters which integrates to infinity.
- Often (but not always!) the resulting "posterior" will be proper.
- It is important that the "posterior" be proper, since otherwise the whole Bayesian framework breaks down.

Pareto example

Recall

$$p(x|\alpha, c) = \frac{\alpha c^\alpha}{x^{\alpha+1}} \mathbb{1}(x > c) \tag{2}$$

$$\mathbb{1}(\alpha, c > 0) \tag{3}$$

Let's derive the posterior:

$$\begin{aligned} p(\alpha, c|x_{1:n}) &\overset{\text{def}}{\underset{\alpha,c}{\propto}} p(x_{1:n}|\alpha, c)p(\alpha, c) \\ &\underset{\alpha,c}{\propto} \mathbb{1}(\alpha, c > 0) \prod_{i=1}^{n} \frac{\alpha c^\alpha}{x_i^{\alpha+1}} \mathbb{1}(x_i > c) \\ &= \frac{\alpha^n c^{n\alpha}}{(\prod x_i)^{\alpha+1}} \mathbb{1}(c < x_*)\mathbb{1}(\alpha, c > 0) \end{aligned} \tag{4}$$

where $x_* = \min\{x_1, \ldots, x_n\}$.

Pareto example

As a joint distribution on (α, c),

- this does not seem to have a recognizable form,
- and it is not clear how we might sample from it directly.

Gibbs sampling

Let's try Gibbs sampling! To use Gibbs, we need to be able to sample $\alpha|c, x_{1:n}$ and $c|\alpha, x_{1:n}$.

By Equation 4, we find that

$$\begin{aligned} p(\alpha|c, x_{1:n}) &\underset{\alpha}{\propto} p(\alpha, c|x_{1:n}) \underset{\alpha}{\propto} \frac{\alpha^n c^{n\alpha}}{(\prod x_i)^\alpha} \mathbb{1}(\alpha > 0) \\ &= \alpha^n \exp\big(-\alpha(\textstyle\sum \log x_i - n \log c)\big) \mathbb{1}(\alpha > 0) \\ &\underset{\alpha}{\propto} \text{Gamma}\left(\alpha \,\middle|\, n+1, \textstyle\sum \log x_i - n \log c\right), \end{aligned}$$

and

$$p(c|\alpha, x_{1:n}) \underset{c}{\propto} p(\alpha, c|x_{1:n}) \underset{c}{\propto} c^{n\alpha} \mathbb{1}(0 < c < x_*),$$

which we will define to be $\text{Mono}(\alpha, x_*)$

Mono distribution

For $a > 0$ and $b > 0$, define the distribution Mono(a, b) (for monomial) with p.d.f.

$$\text{Mono}(x|a, b) \propto x^{a-1}\mathbb{1}(0 < x < b).$$

Since $\int_0^b x^{a-1}dx = b^a/a$, we have

$$\text{Mono}(x|a, b) = \frac{a}{b^a}x^{a-1}\mathbb{1}(0 < x < b),$$

and for $0 < x < b$, the c.d.f. is

$$F(x|a, b) = \int_0^x \text{Mono}(y|a, b)dy = \frac{a}{b^a}\frac{x^a}{a} = \frac{x^a}{b^a}.$$

Pareto example

To use the inverse c.d.f. technique, we solve for the inverse of F on $0 < x < b$: Let $u = \frac{x^a}{b^a}$ and solve for x.

$$u = \frac{x^a}{b^a} \tag{5}$$

$$b^a u = x^a \tag{6}$$

$$b u^{1/a} = x \tag{7}$$

Can sample from Mono(a, b) by drawing $U \sim$ Uniform$(0, 1)$ and setting $X = bU^{1/a}$.[4]

[4] It turns out that this is an inverse of the Pareto distribution, in the sense that if $X \sim$ Pareto(α, c) then $1/X \sim$ Mono$(\alpha, 1/c)$.

Pareto example

So, in order to use the Gibbs sampling algorithm to sample from the posterior $p(\alpha, c|x_{1:n})$, we initialize α and c, and then alternately update them by sampling:

$$\begin{aligned} \alpha|c, x_{1:n} &\sim \text{Gamma}\left(n+1, \sum \log x_i - n \log c\right) \\ c|\alpha, x_{1:n} &\sim \text{Mono}(n\alpha + 1, x_*). \end{aligned}$$

Traceplots

Traceplots. A traceplot simply shows the sequence of samples, for instance $\alpha_1, \ldots, \alpha_N$, or $c_1, \ldots, c_N$. Traceplots are a simple but very useful way to visualize how the sampler is behaving.

Traceplots

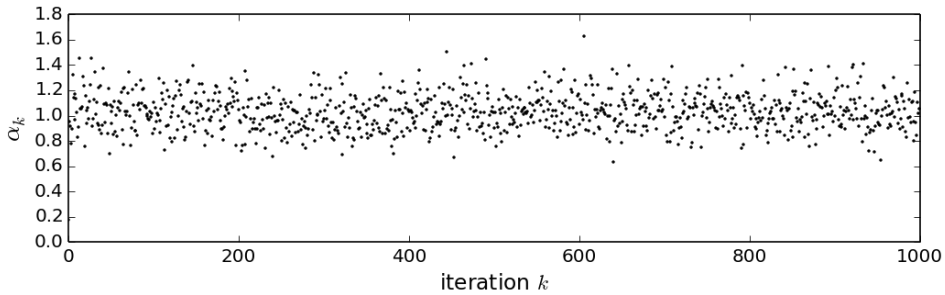

Figure 2: Traceplot of α

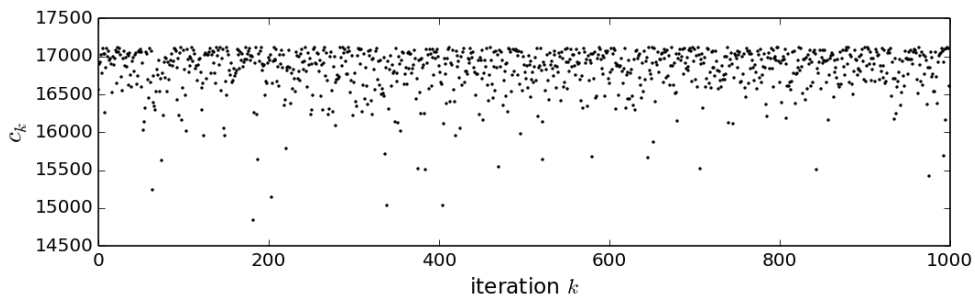

Figure 3: Traceplot of c.

Estimated density

Estimated density. We are primarily interested in the posterior on α, since it tells us the scaling relationship between the size of cities and their probability of occurring.

By making a histogram of the samples $\alpha_1, \ldots, \alpha_N$, we can estimate the posterior density $p(\alpha|x_{1:n})$.

The two vertical lines indicate the lower ℓ and upper u boundaries of an (approximate) 90% credible interval $[\ell, u]$—that is, an interval that contains 90% of the posterior probability:

$$\mathbb{P}(\boldsymbol{\alpha} \in [\ell, u] | x_{1:n}) = 0.9.$$

Estimated density

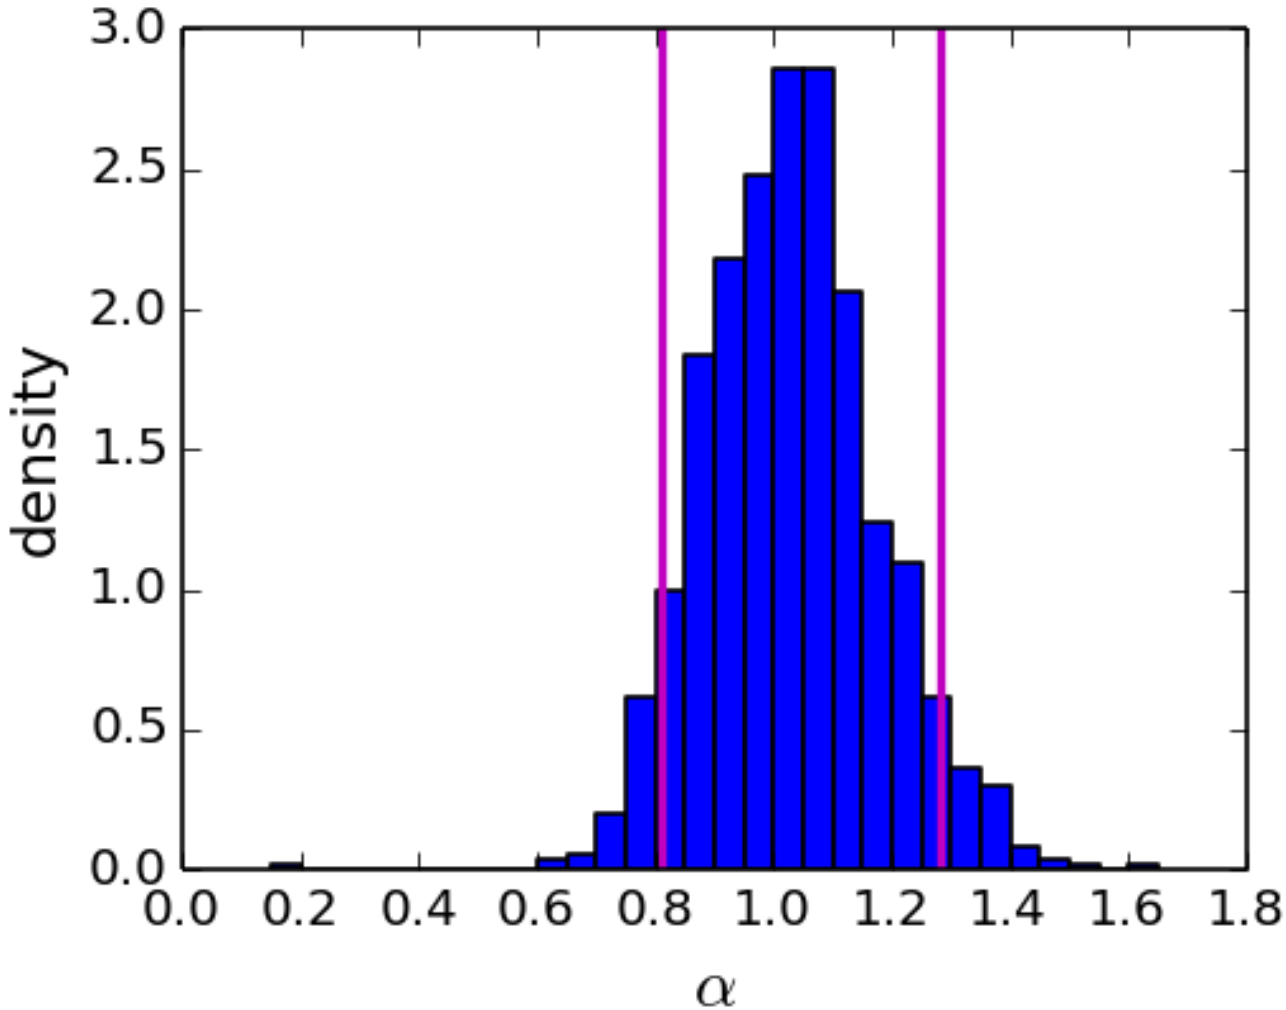

Figure 4: Estimated density of $\alpha|x_{1:n}$ with $\approx$ 90 percent credible intervals.

Running averages

Running averages. Panel (d) shows the running average $\frac{1}{k}\sum_{i=1}^{k}\alpha_i$ for $k = 1, \ldots, N$.

In addition to traceplots, running averages such as this are a useful heuristic for visually assessing the convergence of the Markov chain.

The running average shown in this example still seems to be meandering about a bit, suggesting that the sampler needs to be run longer (but this would depend on the level of accuracy desired).

Running averages

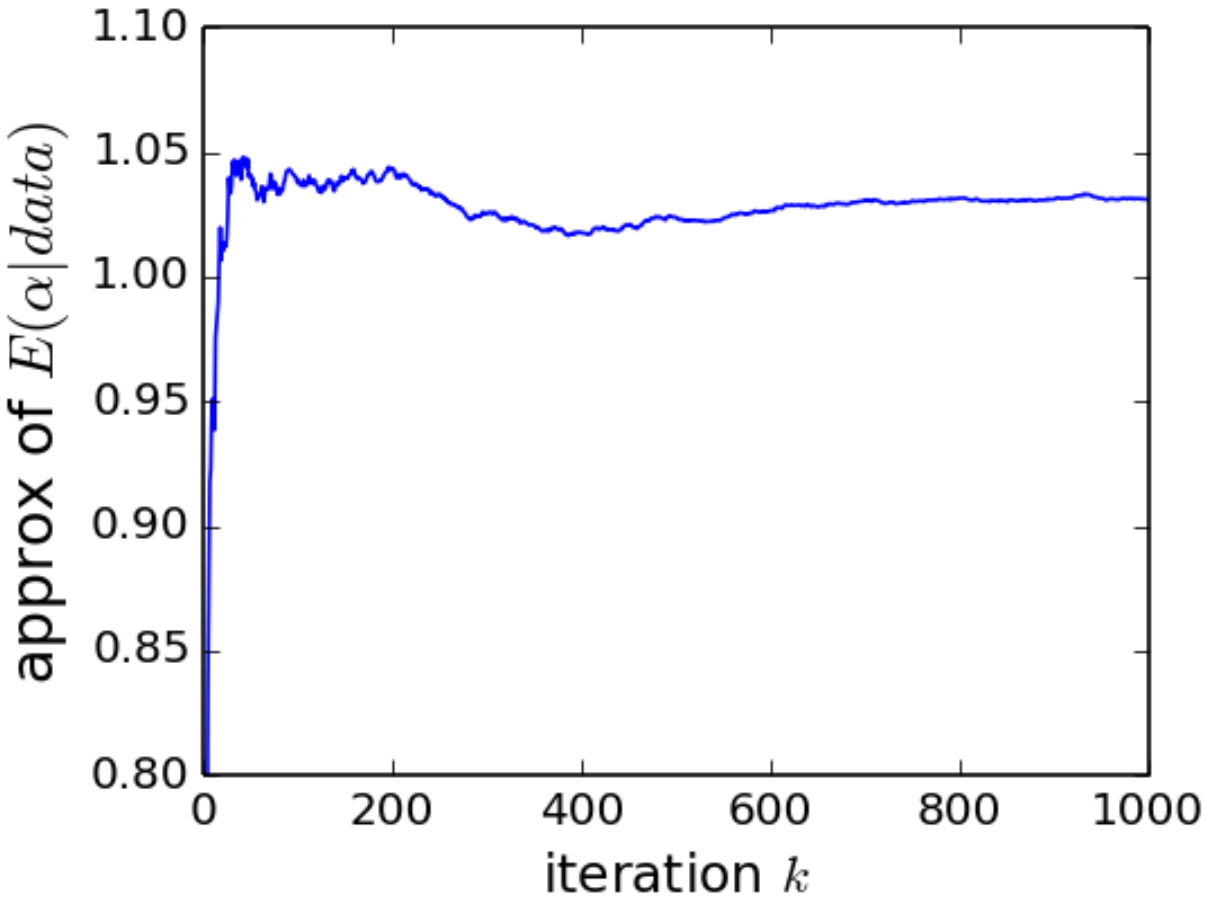

Figure 5: Running average plot

Survival functions

A survival function is defined to be

$$S(x) = \mathbb{P}(X > x) = 1 - \mathbb{P}(X \leq x).$$

Power law distributions are often displayed by plotting their survival function $S(x)$, on a log-log plot.

Why? $S(x) = (c/x)^{\alpha}$ for the Pareto(α, c) distribution and on a log-log plot this appears as a line with slope $-\alpha$.

The posterior survival function (or more precisely, the posterior predictive survival function), is $S(x|x_{1:n}) = \mathbb{P}(X_{n+1} > x \mid x_{1:n})$.

Survival functions

Figure 6(e) shows an empirical estimate of the survival function (based on the empirical c.d.f., $\hat{F}(x) = \frac{1}{n}\sum_{i=1}^{n} \mathbb{1}(x \geq x_i)$) along with the posterior survival function, approximated by

$$S(x|x_{1:n}) = \mathbb{P}(X_{n+1} > x \mid x_{1:n}) \tag{8}$$

$$= \int \mathbb{P}(X_{n+1} > x \mid \alpha, c) p(\alpha, c|x_{1:n}) d\alpha dc \tag{9}$$

$$\approx \frac{1}{N}\sum_{i=1}^{N} \mathbb{P}(X_{n+1} > x \mid \alpha_i, c_i) = \frac{1}{N}\sum_{i=1}^{N} (c_i/x)^{\alpha_i}. \tag{10}$$

This is computed for each x in a grid of values.

Survival functions

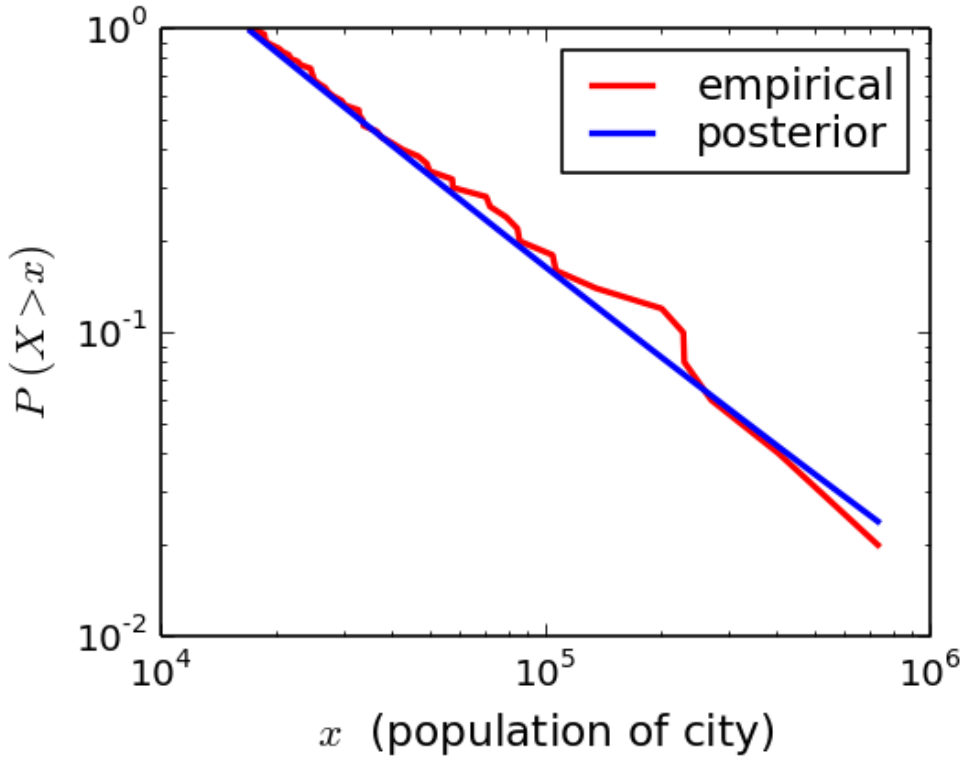

Figure 6: Empirical vs posterior survival function