

# Module 12: Linear Regression, the g-prior, and model selection

Rebecca C. Steorts

# Agenda

# Setup

- ▶  $X_{n \times p}$ : regression features or covariates (design matrix)
- ▶  $x_{p \times 1}$ :  $i$ th row vector of the regression covariates
- ▶  $y_{n \times 1}$ : response variable (vector)
- ▶  $\beta_{p \times 1}$ : vector of regression coefficients

Goal: Estimation of  $p(y \mid x)$ .

Dimensions:  $y_i - \beta^T x_i = (1 \times 1) - (1 \times p)(p \times 1) = (1 \times 1)$ .

## Multivariate Setup

Let's assume that we have data points  $(x_i, y_i)$  available for all  $i = 1, \dots, n$ .

- ▶  $y$  is the response variable

$$y = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}_{n \times 1}$$

- ▶  $x_i$  is the  $i$ th row of the design matrix  $X_{n \times p}$ .

Consider the regression coefficients

$$\beta = \begin{pmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{pmatrix}_{p \times 1}$$

## Multivariate Setup

$$y \mid X, \beta, \sigma^2 \sim MVN(X\beta, \sigma^2 I)$$

$$\beta \sim MVN(\beta_0, \Sigma_0)$$

Recall the posterior can be shown to be

$$\beta \mid \mathbf{y}, \mathbf{X} \sim MVN(\beta_n, \Sigma_n)$$

where

$$\beta_n = E[\beta \mid \mathbf{y}, \mathbf{X}, \sigma^2] = (\Sigma_o^{-1} + (X^T X)^{-1}/\sigma^2)^{-1}(\Sigma_o^{-1}\beta_0 + \mathbf{X}^T \mathbf{y}/\sigma^2)$$

$$\Sigma_n = \text{Var}[\beta \mid \mathbf{y}, \mathbf{X}, \sigma^2] = (\Sigma_o^{-1} + (X^T X)^{-1}/\sigma^2)^{-1}$$

How do we specify  $\beta_0$  and  $\Sigma_0$ ?

## The g-prior

To do the *least amount of calculus*, we can put a *g-prior* on  $\beta$

$$\beta \mid \mathbf{X}, \mathbf{z} \sim \text{MVN}(0, g \sigma^2 (\mathbf{X}^T \mathbf{X})^{-1}).$$

$$\beta_n = E[\beta \mid \mathbf{y}, \mathbf{X}, \sigma^2] = \frac{g}{g+1} (\Sigma_o^{-1} + (X^T X)^{-1} / \sigma^2)^{-1} = \frac{g}{g+1} \hat{\beta}_{ols}$$

$$\Sigma_n = \text{Var}[\beta \mid \mathbf{y}, \mathbf{X}, \sigma^2] = \frac{g}{g+1} (X^T X)^{-1} / \sigma^2 = \frac{g}{g+1} \text{Var}[\hat{\beta}_{ols}]$$

- ▶  $g$  shrinks the coefficients and can prevent overfitting to the data
- ▶ if  $g = n$ , then as  $n$  increases, inference approximates that using  $\hat{\beta}_{ols}$

## Variance component $\sigma^2$

What about a prior on  $1/\sigma^2 = \lambda$

$$y \mid X, \beta, \sigma^2 \sim MVN(X\beta, \sigma^2 I) \quad (1)$$

$$\lambda \sim \text{Gamma}(\nu_0/2, \nu_0\lambda^{-1}/2) \quad (2)$$

Then the posterior can be shown to be

$$p(\lambda \mid y, X) \sim \text{Gamma}([\nu_0 + n]/2, [\nu_0\lambda^{-1} + SSR_g]/2)$$

where  $SSR_g$  is somewhat complicated (see Hoff for details, p. 158).

## Variance component $\sigma^2$

The joint distribution can be written as

$$p(\lambda^{-1}, \beta \mid y, X) = p(\lambda^{-1} \mid y, X) \times p(\beta \mid y, X, \lambda^{-1})$$

Goal: simulate  $(\lambda^{-1}, \beta) \sim p(\lambda^{-1}, \beta \mid y, X)$  Starting value  $(\beta_0, \lambda_0)$

1. Simulate

$$\lambda^{-1} \sim p(\lambda^{-1} \mid y, X)$$

Gives us  $(\lambda_1^{-1}, \beta_0)$  2. Use this updated value of  $\lambda_1^{-1}$  to simulate

$$\beta \sim p(\beta \mid y, X, \lambda^{-1})$$

Gives us  $(\lambda_1^{-1}, \beta_1)$

Run the sampler for  $S$  iterations.



## Back to the oxygen uptake example

$$y \mid X, \beta, \sigma^2 \sim \text{MVN}(X\beta, \lambda^{-1}I) \quad (3)$$

$$\beta \mid \lambda \text{MVN}(0, g(X^T X)^{-1}) \quad (4)$$

$$\lambda^{-1} \sim \text{Gamma}(\nu_0/2, \nu_0 \lambda^{-1}/2) \quad (5)$$

We will use the g-prior, where

1.  $g = n$
2.  $\nu_0 = 1$
3.  $\sigma_0^2 = \lambda^{-1} = \hat{\sigma}_{ols} = 8.54$

## Application to diabetes (Exercise 9.2, part a)

As described in Exercise 7.6, suppose we have data on health-related variables of a population of 532 women.

In this exercise we will be modeling the conditional distribution of glucose level ( $\text{glu}$ ) as a linear combination of the other variables, excluding the variable diabetes.

## Regression model on the g-prior

Fit a regression model using the g-prior with  $g = n$ ,  $\nu_0 = 2$  and  $\sigma_0^2 = 1$ . Obtain posterior confidence intervals for all of the parameters.

## Regression model on the g-prior

Section 9.2.2 (Hoff) shows that under the  $g$  prior,  $p(\sigma^2 \mid \mathbf{y}, \mathbf{X})$  and  $p(\boldsymbol{\beta} \mid \mathbf{y}, \mathbf{X}, \sigma^2)$  are inverse gamma and multivariate normal distributions respectively.

## Regression model on the g-prior

Therefore samples from the joint posterior  $p(\sigma^2, \beta \mid \mathbf{y}, \mathbf{X}, \sigma^2)$  can be made with a Monte Carlo approximation.

We first center and scale all the variables so that there is no need to include an intercept in the model.

## Regression model on the g-prior

```
library(knitr)
rm(list=ls())
azd_data = read.table("azdiabetes.dat", header = TRUE)
head(azd_data)
```

##	npreg	glu	bp	skin	bmi	ped	age	diabetes
## 1	5	86	68	28	30.2	0.364	24	No
## 2	7	195	70	33	25.1	0.163	55	Yes
## 3	5	77	82	41	35.8	0.156	35	No
## 4	0	165	76	43	47.9	0.259	26	No
## 5	0	107	60	25	26.4	0.133	23	No
## 6	5	97	76	27	35.6	0.378	52	Yes

## Regression model on the g-prior

```
y = azd_data$glu  
X = as.matrix(azd_data[,c(-2,-8)])  
head(X)
```

##		npreg	bp	skin	bmi	ped	age
##	[1,]	5	68	28	30.2	0.364	24
##	[2,]	7	70	33	25.1	0.163	55
##	[3,]	5	82	41	35.8	0.156	35
##	[4,]	0	76	43	47.9	0.259	26
##	[5,]	0	60	25	26.4	0.133	23
##	[6,]	5	76	27	35.6	0.378	52

# Standardization

```
# standardize data to have mean 0 and variance 1  
ys = scale(y)  
Xs = scale(X)  
n = dim(Xs)[1]  
p = dim(Xs)[2]
```



# Hyper-parameters

## hyper-parameters

```
g = n  
nu0 = 2  
s20 = 1
```

## Intermediate Matrices

```
# intermediate matrices
```

```
Hg = (g/(g+1)) * Xs %*% solve(t(Xs) %*% Xs) %*% t(Xs)
```

```
SSRg = t(ys) %*% ( diag(1,nrow=n) - Hg ) %*% ys
```

## Monte carlo

```
# number of posterior samples
S = 1000

# generate posteriors
s2 = 1/rgamma(S, (nu0+n)/2, (nu0*s20 + SSRg)/2)
Vb = g * solve(t(Xs) %*% Xs)/(g+1)
Eb = Vb %*% t(Xs) %*% ys
E = matrix(rnorm(S*p, 0, sqrt(s2)),S,p)
beta_s = t( t(E %*% chol(Vb)) + c(Eb))

# transform coefficients to the original scale
sd_X = apply(X,2,sd)
Beta_a = sweep(beta_s,2,sd_X,FUN = "/" )
```

## The 95% posterior confidence intervals

```
# 95% credible interval
```

```
Beta_CIA = apply(Beta_a, 2, quantile, c(0.025, 0.975))  
kable(data.frame(Beta_CIA))
```

	npreg	bp	skin	bmi	ped
2.5%	-0.0515254	-0.0005462	-0.0033335	0.0059186	0.1047847
97.5%	0.0091583	0.0133459	0.0159199	0.0352423	0.5617721

# Model selection

- ▶ Often we have a large number of covariates.
- ▶ Using all of them induces poor statistical performance.
- ▶ How can we reduce the covariates and have good inference and prediction?
- ▶ Common method: Backwards and stepwise regression (slow).

## Model selection

Suppose that we believe some of the regression coefficients are 0.

Come up with a prior distribution that reflects the probability of this occurring.

Consider

$$y_i = z_1 b_1 x_{i,1} + \dots z_p b_p x_{i,p},$$

where  $b_p$  is a real number and  $z_j$  indicate which regression coefficients are nonzero.

Note:  $\beta_j = b_j \times z_j$ .

# Bayesian model selection

Bayesian model selection works by obtaining a posterior distribution for  $\mathbf{z}$ .

Assume a prior  $p(\mathbf{z})$ .

Then

$$p(\mathbf{z} \mid \mathbf{Y}, \mathbf{X}) = \frac{p(\mathbf{z})p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z})}{\sum_{\mathbf{z}} p(\mathbf{z})p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z})}$$



## Bayesian model selection

Suppose we want to compare two models  $z_a$  and  $z_b$ . Consider

$$\text{odds}(z_a, z_b \mid \mathbf{Y}, \mathbf{X}) = \frac{p(z_a \mid \mathbf{Y}, \mathbf{X})}{p(z_b \mid \mathbf{Y}, \mathbf{X})} = \frac{p(z_a)}{p(z_b)} \times \frac{p(\mathbf{Y} \mid \mathbf{X}, z_a)}{p(\mathbf{Y} \mid \mathbf{X}, z_b)}$$

This is posterior odds = prior odds  $\times$  "Bayes factor"

"Bayes factor": how much the data favor model  $z_a$  over model  $z_b$

To obtain a posterior distribution over models, we must compute  $p(\mathbf{Y} \mid \mathbf{X}, z)$  for *each* model under consideration.

## Bayesian model selection

We must compute

$$p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z}) = \int \int p(\mathbf{Y}, \beta, \sigma^2, \mid \mathbf{X}, \mathbf{z}) \quad (6)$$

$$\int \int p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z}) p(\beta \mid \mathbf{X}, \mathbf{z}) p(\sigma^2). \quad (7)$$

To do the *least amount of calculus*, we can put a *g-prior* on  $\beta$

$$\beta \mid \mathbf{X}, \mathbf{z} \sim \text{MVN}(0, g \sigma^2 (\mathbf{X}^T \mathbf{X})^{-1}).$$

## Back to the g-prior

Given the g-prior

$$\boldsymbol{\beta} \mid \mathbf{X}, \mathbf{z} \sim MVN(0, g \sigma^2 (\mathbf{X}^T \mathbf{X})^{-1}),$$

$p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z})$  can be worked out in closed form (details p. 165).

Go through the details on your own.

## Back to the g-prior

This results in being able to compute

$$\frac{p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z}_a)}{p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z}_b)} = (1 + n)^{(p_{z_b} - p_{z_a})/2} \times \left( \frac{s_{z_a}^2}{s_{z_b}^2} \right)^{1/2} \quad (8)$$

$$\times \left( \frac{s_{z_b}^2 + SSR_g^{z_b}}{s_{z_b}^2 + SSR_g^{z_a}} \right)^{(n+1)/2} \quad (9)$$

We have a ratio of the marginal probabilities, giving us a balance between model complexity and model fit.

Suppose  $p_{z_b}$  is large compared to  $p_{z_a}$ .

This causes a penalization of model  $z_b$

Note that a large value of  $SSR_g^{z_b}$  compared to  $SSR_g^{z_a}$  will penalize model  $z_a$ .

# Bayesian Model Averaging

Suppose that we are content with a estimate of  $\beta$  from which we can make predictions.

We may also want a list of relatively high probablitiy models.

We can use a Markov chain to search through the space of models for values of  $z$  with high posterior probability.

## Bayesian model averaging

Suppose  $p$  is large. Then  $2^p$  models to consider.

Instead let's use a Gibbs sampler to search through the space of models for values where  $\mathbf{z}$  has a high posterior probability.

Generate a new value of  $\mathbf{z}$  via

$$p(z_j \mid \mathbf{Y}, \mathbf{X}, \mathbf{z}_{-j}).$$

The full conditional that  $z_j = 1$  can be written as  $o_j / (o_j + 1)$ .

$$o_j = \frac{p(z_j = 1 \mid \mathbf{Y}, \mathbf{X}, \mathbf{z}_{-j})}{p(z_j = 0 \mid \mathbf{Y}, \mathbf{X}, \mathbf{z}_{-j})} \quad (10)$$

$$= \frac{p(z_j = 1)p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z}_{-j}, z_j = 1)}{p(z_j = 0)p(\mathbf{Y} \mid \mathbf{X}, \mathbf{z}_{-j}, z_j = 0)} \quad (11)$$

## Bayesian model averaging

Note: we may also want to obtain posterior samples of  $\beta$  and  $\sigma^2$ .

Using the conditional distributions from Section 9.2, we can sample from these directly.

The Gibbs sampling scheme requires using Section 9.2 and 9.3 (covered in lab).

# Bayesian model averaging

$$\begin{array}{ccccc} \mathbf{z}^{(s)} & \longrightarrow & \sigma^{2(s)} & \longrightarrow & \boldsymbol{\beta}^{(s)} \\ \downarrow & & & & \\ \mathbf{z}^{(s+1)} & \longrightarrow & \sigma^{2(s+1)} & \longrightarrow & \boldsymbol{\beta}^{(s+1)} \end{array}$$

Figure 1: Start with  $\mathbf{z}^{(s)}$ . Then in random order update  $z_j$  from its full conditional.



# Bayesian model averaging

Generate

$$\{\mathbf{z}^{(s+1)}, \sigma^{2(s+1)}, \beta^{(s+1)}\} :$$

1. Set  $\mathbf{z} = \mathbf{z}^{(s)}$
2. For  $j \in \{1, \dots, p\}$  in random order, replace  $z_j$  with a sample from

$$p(z_j \mid \mathbf{z}_{-j}, \mathbf{Y}, \mathbf{X})$$

3. Set  $\mathbf{z}^{(s+1)} = \mathbf{z}$
4. Sample  $\sigma^{2(s+1)} \sim p(\sigma^2 \mid \mathbf{z}^{(s+1)}, \mathbf{Y}, \mathbf{X})$
5. Sample  $\beta^{(s+1)} \sim p(\beta \mid \mathbf{z}^{(s+1)}, \sigma^{2(s+1)}, \mathbf{Y}, \mathbf{X})$

## Back to diabetes data (Exercise 9.2, b)

Let's perform Bayesian model averaging (as described in Section 9.3)

Obtain  $P(\beta_j \neq 0 | y)$  as well as posterior confidence intervals for all of the parameters. Compare our results to that in part (a.)

## Back to diabetes data (Exercise 9.2, b)

The following function `lpy.X` calculates the log of  $p(\mathbf{y} \mid \mathbf{X})$ , which we will use in implementing the Gibbs sampler for Bayesian model averaging.

```
## a function to compute the marginal probability
lpy.X = function(y, X, g=length(y), nu0=1, s20=try(summary
  n = dim(X)[1]
  p = dim(X)[2]
  if (p==0) { Hg = 0; s20 = mean(y^2)}
  if (p>0){ Hg = (g/(g+1)) * X %>% solve(t(X) %>% X) %>% t
  SSRg = t(y) %>% ( diag(1, nrow=n) - Hg ) %>% y -
    .5*( n*log(pi) + p*log(1+g) + (nu0+n)*log(nu0*s20 + SSR
    lgamma( (nu0+n)/2 ) - lgamma(nu0/2)
}
```

## Back to diabetes data (Exercise 9.2, b)

Let  $\mathbf{z}$  be the random binary vector of variable indicators. Generating samples of  $p(\mathbf{z}, \sigma^2, \beta)$  from the joint posterior distribution is achieved with the following steps:

1. For  $j \in \{1, \dots, p\}$  in random order, draw  $z_j$  from  $p(z_j \mid \mathbf{z}_{-j}, \mathbf{y}, \mathbf{X})$ .
2. Sample  $\sigma^2 \sim p(\sigma^2 \mid \mathbf{z}, \mathbf{y}, \mathbf{X})$ .
3. Sample  $\beta \sim p(\beta \mid \mathbf{z}, \sigma^2, \mathbf{y}, \mathbf{X})$ .

# Gibbs sampler

```
## MCMC setup
g = n
nu0 = 1 # unit information prior
z = rep(1, p)
lpy.c = lpy.X(ys, Xs[,z==1,drop=FALSE])
S = 10000
Z = matrix(NA, S, p)
# Sigma2 = numeric(S)
B = matrix(0, S, p)

## Gibbs sampler
for(s in 1:S){
  # if(s %% 100 ==0) {print(s)}
  # sample z
  for (j in sample(1:p)){
    zp = z
    zp[j] = 1 - zp[j]
    lpy.p = lpy.X(ys,Xs[, zp==1, drop=FALSE])
```